

QDMN: Quad-Directional Modular Networks as Explicit Multi-Route Processing Fabrics

Sohan Poudel
forreserachonlyso@gmail.com
ORCID: 0009-0004-5762-3898

April 26, 2026

Transparency / provenance note

Origin of the idea. The initial motivating question for this draft was: *can “Up” and “Down” directions be treated as first-class processing directions in neural networks (in addition to the usual forward/backward framing)?*

AI-assisted drafting. Portions of the text, code scaffolding, and experimental write-up in this document were produced with the assistance of large language models (LLMs) (e.g., ChatGPT and Gemini) and then curated/edited by the author.

Conversation and public record. For transparency, the author documents (i) the originating chat conversations and (ii) an earlier public post announcing the idea. (Identifiers are listed as plain text; replace/update as desired.)

- Originating chat (ChatGPT): chatgpt.com/share/6910da26-9cc4-8012-bd6e-6aab73040a68
- Originating chat (Gemini): gemini.google.com/share/263e0e5409ab
- Chat used for this paper (ChatGPT): chatgpt.com/share/697b7ce7-1f40-8012-bdd7-b83ede452531
- Chat used for this paper (Gemini): gemini.google.com/share/54337111c4cd
- Public post (tweet): twitter.com/thgie_8_eight/status/1987591129718820994
- OLD Colab (first version): colab.research.google.com/drive/1MrISPV4QtIAyn6SSe1X50xsaN40g1h0E?usp=sharing
- New Colab (paper version): colab.research.google.com/drive/134Kns6B3awWhS_i-LA5QYFyncaAw6K_W?usp=sharing

Disclaimer. This document is a research draft and should not be interpreted as a peer-reviewed claim of novelty or performance; readers are encouraged to verify experiments and prior art independently.

Abstract

We introduce *Quad-Directional Modular Networks* (QDMN), an architecture-first framework that treats computation as an explicit multi-directional fabric rather than a single feed-forward pipeline. QDMN exposes four semantic processing directions—Forward (F), Backward (B), Up (U), and Down (D)—implemented as registrable modules with declared input/output feature dimensions. A *route* (e.g., [F, U, D]) deterministically advances a central *CoreNode* state through compatible modules, enabling hand-designed or algorithmically selected multi-stage inference and reconstruction cycles.

Contributions.

- A routing-and-registry abstraction for multi-directional module composition with dimension-aware dispatch.
- A minimal PyTorch reference implementation (Appendix) suitable as a reproducibility artifact.
- Suggested evaluation protocols for U-to-D reconstruction, F-to-B refinement/consistency, and route-length adaptive compute.

1 Introduction

Neural networks are commonly described in terms of a forward pass (computing predictions) and a backward pass (computing gradients). However, these notions conflate two distinct ideas: (i) the *semantic direction* of information processing (e.g., refining, abstracting, reconstructing) and (ii) the *mathematical direction* used by training algorithms (e.g., reverse-mode automatic differentiation).

This paper explores a framework question: can we make multiple *semantic* processing directions first-class primitives, so that inference is an explicit, programmable composition of directional transformations?

We propose **QDMN** (*Quad-Directional Modular Network*), a routing framework that:

- assigns modules to four direction symbols: **F** (forward/progression), **B** (backward/refinement or consistency), **U** (up/compression or abstraction), and **D** (down/expansion or reconstruction);
- stores a mutable representation in a central *CoreNode*;
- executes explicit *routes*—sequences of direction symbols—with dimension-aware module dispatch.

What is new here (in this draft). The individual ideas of “up/down” (encoder–decoder processing) and iterative refinement exist in many architectures. The intended novelty of QDMN is the *interface*: a general, compositional *fabric* in which multiple directional modules can be registered and recombined. In particular:

- **Up/Down as reusable primitives:** “U” and “D” are not tied to one fixed encoder–decoder graph; they are modules that can appear multiple times and interleave with refinement steps.
- **Route-as-program:** inference is controlled by an explicit route (e.g., [F, U, D, B, H]), enabling route switching at test time.
- **Dimension-typed registry:** modules are keyed by input feature dimension, making compatibility checks and multi-scale routing explicit.

- **Toward grid-like composition:** the same abstractions can be extended from 1-D routes to 2-D “processing fabrics” (a grid of reusable directional transitions), although this draft focuses on sequential routes.

QDMN is intentionally *training-agnostic*: it specifies a compositional scaffold rather than a particular loss, optimizer, or a learned router.

Route length as a controllable variable. Unlike fixed feed-forward pipelines, QDMN exposes inference depth via route length. Applying additional refinement steps (e.g., appending extra B modules) is a form of test-time compute scaling, but monotonic improvement with depth is not guaranteed without route-aware training objectives or explicit stopping criteria.

2 Related Work (Positioning)

QDMN is most closely related in spirit to (i) encoder–decoder and multi-scale architectures (where “up” and “down” correspond to compression/expansion), (ii) iterative refinement and feedback-style networks, and (iii) modular composition and routing/conditional-computation frameworks.

This draft does *not* claim to invent compression/expansion or feedback. Instead, QDMN aims to provide a small, explicit *programming model* for composing these behaviors: a dimension-aware registry plus a route execution semantics. Unlike conventional encoder–decoder systems where the up/down path is baked into a single graph, QDMN treats U and D as reusable modules that can be invoked, repeated, and interleaved with refinement steps; and unlike learned routers (e.g., mixture-of-experts variants), QDMN foregrounds explicit, human-specified routes as a controllable diagnostic and compute knob.

A complete citation-backed comparison to prior work is left for a full version.

3 QDMN Framework

3.1 Core Abstractions

Let $x \in \mathbb{R}^{\cdots \times d}$ denote a tensor whose last dimension is the *feature dimension* d . QDMN maintains a mutable state s in a **CoreNode**, where $\dim(s) = d$.

A **directional module** is a function (implemented as an `nn.Module`)

$$m_{\sigma,d} : \mathbb{R}^{\cdots \times d} \rightarrow \mathbb{R}^{\cdots \times d'}$$

parameterized by a *direction symbol* $\sigma \in \{F, B, U, D\}$ and an input feature dimension d . Each module declares an output feature dimension d' .

QDMN maintains a registry keyed by input dimension, mapping direction symbols to modules:

$$\mathcal{R}[d][\sigma] \mapsto (m_{\sigma,d}, d').$$

3.2 Routing Semantics

A **route** is a sequence of direction symbols

$$\pi = [\sigma_1, \dots, \sigma_T].$$

Given an initial tensor x with feature dimension d_0 , QDMN executes:

$$\begin{aligned} s_0 &= x \\ s_t &= m_{\sigma_t, \dim(s_{t-1})}(s_{t-1}) \quad \text{for } t = 1, \dots, T. \end{aligned}$$

Execution fails if a required module is not registered for the current feature dimension.

Interpretation of directions. The direction symbols are semantic labels rather than mathematical inverses:

- **F (Forward):** advance a representation (prediction, progression).
- **B (Backward):** refine, critique, or check consistency against constraints.
- **U (Up):** compress/abstract (reduce feature dimension or complexity).
- **D (Down):** expand/reconstruct (increase detail or feature dimension).

3.3 Design Notes

Dimensional contracts. QDMN uses the last dimension of tensors as the routing “type”. This supports vector, sequence, and image data, provided modules preserve all non-feature dimensions (or do so intentionally).

Adapters. To connect incompatible feature dimensions across routes or across multiple QDMNs, adapters (e.g., linear projections) can be inserted.

Training. QDMN does not prescribe training. Common options include (i) route-conditioned supervised loss, (ii) reconstruction losses for [U, D], and (iii) consistency losses for [F, B].

4 Reference Implementation

We provide a minimal PyTorch implementation that supports: (i) a **CoreNode** that stores the current state, (ii) a registry of modules keyed by input feature dimension and direction symbol, and (iii) an explicit **route(path, x)** method that executes the route.

4.1 Complexity

Let T be route length and let each module cost C_t . Route execution costs $\sum_{t=1}^T C_t$ with $\mathcal{O}(1)$ registry lookup per step.

5 Experiments

5.1 Goals

We study QDMN primarily as an architectural abstraction rather than as a single task-specific model. Our experiments target three questions: (i) can QDMN express common computation motifs (direct prediction, compress–reconstruct, iterative refinement), (ii) can QDMN match simple baselines on toy tasks under comparable training recipes, and (iii) what happens when we execute *longer* routes than those used during training (route switching / depth scaling).

Table 1: Benchmark summary (single-seed run). Replace values with the outputs produced by your current run and report mean $\pm$ std over multiple seeds in a final version.

Task	Metric	Baseline 1	Baseline 2	QDMN route	Value
8 $\times$ 8 image cls.	Accuracy	CNN	MLP	[F, H]	1.00
8 $\times$ 8 image cls.	Accuracy	CNN	MLP	[U, D, H]	1.00
Sine regression	MSE $\downarrow$	MLP	—	[F, B, D]	0.013
Seq. parity (len=10)	Accuracy	MLP	RNN	[F, B, D, H]	0.60

Table 2: Depth generalization on sine regression. Naive depth-1 training exhibits deterioration at larger depths (e.g., depth 8), while multi-depth training keeps performance bounded across depths in this run.

Depth d	QDMN (naive) MSE $\downarrow$	QDMN (multi-depth) MSE $\downarrow$
1	0.121418	0.169901
2	0.104954	0.143478
4	0.154913	0.112439
8	0.658170	0.216256

5.2 Tasks

We use three small synthetic tasks (chosen for fast iteration and clear diagnostics): (1) 8 $\times$ 8 image classification, (2) 1-D noisy sine regression, and (3) sequence parity (length 10). Full data generation and model definitions are provided in Appendix A.

5.3 Training setup

Unless noted otherwise, models are trained with Adam (lr = 10^{-3}), batch size 64, and a fixed random seed. We report accuracy for classification and mean-squared error (MSE) for regression.

6 Results and Discussion

6.1 Benchmark summary

Table 1 reports representative results from our reference script. These experiments are not intended to establish state-of-the-art performance; they serve to validate that QDMN routes can be trained end-to-end and to highlight route-dependent behavior.

6.2 Depth generalization: naive vs stable refiners

A central motivation for QDMN is that inference depth is programmable via route length. However, extending a route beyond its training depth is not guaranteed to improve performance. We therefore evaluate *depth generalization* on a toy sine regression task by comparing: (i) a *naive* QDMN trained only at depth 1, and (ii) a *stable* QDMN trained with a multi-depth objective.

Concretely, the model computes $\hat{y} = D(B^d(\text{ReLU}(F(x))))$ where d is the inference depth (number of applications of B). We use a gated residual refiner with spectral normalization as a practical stability mechanism.




Figure 1: Depth generalization behavior on sine regression (log-scaled depth axis). In this run, the naive model diverges at larger depths, while multi-depth training keeps errors bounded.

Table 3: Multi-seed summary (validation/test MSE; lower is better).

Model/setting	Val. MSE (mean $\pm$ std)	Test MSE (mean $\pm$ std)
MLP (depth 1)	0.014721 $\pm$ 0.002424	0.015340 $\pm$ 0.001809
QDMN_naive (depth 1)	0.018939 $\pm$ 0.002827	0.019220 $\pm$ 0.001577
QDMN_stable (depth 1)	0.107850 $\pm$ 0.046955	0.099328 $\pm$ 0.042365
QDMN_stable (depth 4)	0.072137 $\pm$ 0.035018	0.064939 $\pm$ 0.034157
QDMN_stable (depth 8)	0.066954 $\pm$ 0.010260	0.059372 $\pm$ 0.014283

Interpretation. The results support a conservative claim: *depth generalization is conditional, not automatic*. When a refinement operator B is trained only for a single application, repeated application can move the representation outside the manifold expected by the decoder. Multi-depth training and gated residual refiners mitigate this by explicitly optimizing behavior across multiple unroll depths.

Takeaway. QDMN makes depth scaling a controllable variable and provides a clean testbed for studying stability mechanisms (gating, contraction/regularization, and multi-depth objectives) required for reliable iterative refinement.

6.3 Multi-seed results (paper-ready script)

We additionally ran a larger multi-seed experiment on noisy sine regression using three random seeds (42, 123, 999) and reported mean $\pm$ standard deviation. The full script is included in Appendix D.




Figure 2: Semantic compositionality demo on noisy sine filtering. A deeper route $[F, U, D, B^3, H]$ can apply multiple refinement steps compared to a short route $[F, H]$.

Depth sweep behavior. In this run, QDMN_naive exhibits severe degradation as depth increases (e.g., test MSE grows from ≈ 0.019 at depth 1 to ≈ 5594 at depth 32), while QDMN_stable remains comparatively bounded through moderate depths (e.g., best mean test MSE around depth 8) before degrading at very large depths.

Statistical note. A paired t-test comparing QDMN_stable at depth 4 versus the MLP baseline (test MSE across the three seeds) yielded $p \approx 0.138$ in this run; this small experiment is underpowered and is included primarily to illustrate the evaluation workflow rather than to claim statistical superiority.

6.4 Semantic compositionality: $[F, U, D, B^n, H]$

We include an additional toy experiment that composes multiple semantic directions to perform noise filtering on a noisy sine signal. The model is trained on a specific route $[F, U, D, B, H]$ and then evaluated using different inference-time routes (e.g., a short route $[F, H]$ versus a deeper refinement route $[F, U, D, B^3, H]$). This experiment is intended as a qualitative demonstration of route-based logic rather than a benchmark.

7 Limitations

- QDMN is a compositional framework; performance depends on module choices and training objectives.
- Explicit routing requires external logic (heuristics or a controller) and may be brittle without careful design.

- Reusing the same refinement module multiple times at test time can be out-of-distribution relative to training; without multi-depth training or a stopping rule, additional steps may degrade performance.
- If modules change tensor shapes beyond the last dimension, additional bookkeeping is required.

8 Conclusion

QDMN reframes neural computation as an explicit multi-directional program over a shared core state. By making routing and directionality first-class, QDMN provides a compact scaffold for experimentation with refinement, abstraction, reconstruction, and adaptive inference.

A Reproducibility Artifact (Single-Block Script)

We include a single-file benchmark script (data generation, baselines, QDMN variants, and the route-switching diagnostic) as a reproducibility artifact.

```
"""
QDMN: Framework + Benchmark Single-Block Script
Paste in Colab (or run in any Python env with torch installed).

Purpose:
- Provide a concise, reproducible demonstration of QDMN (Quad-Directional Modular Network)
  as a general framework (register modules, run routes).
- Provide a set of benchmark experiments comparing QDMN to common baselines.
- Include the "Route-Switching / Novelty" test that shows QDMN's runtime route flexibility.

Notes:
- This is a compact demo for a paper. Increase model sizes / dataset sizes / epochs for stronger claims.
- The QDMN framework operates on the *last* dimension of tensors. For structured data (images),
  we flatten or use Conv modules that respect channel/height/width examples provided.

Author: (You) include this cell as the reproducible artifact for your paper.
"""

import time, random, os
from typing import Dict, List, Optional
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

# -----
```

```

# Config (tweak as needed)
# -----
SEED = 42
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", DEVICE)

# Default hyperparameters (paper: report these)
EPOCHS_SMALL = 25 # small for quick runs; increase for stronger claims
EPOCHS_LARGE = 60 # for regression or when you need better fit
BATCH = 64
LR = 1e-3

# -----
# QDMN Framework
# -----
class CoreNode(nn.Module):
    def __init__(self, initial_dim: int):
        super().__init__()
        self._state: Optional[torch.Tensor] = None
        self._dim = initial_dim

    def set(self, x: torch.Tensor):
        if x.shape[-1] != self._dim and self._state is None:
            # allow initial state to set dim if created differently, but prefer consistent
            initial_dim
            self._dim = x.shape[-1]
        self._state = x
        self._dim = x.shape[-1]

    def get(self) -> torch.Tensor:
        if self._state is None:
            raise RuntimeError("CoreNode_empty: call set(x) first")
        return self._state

    @property
    def dim(self) -> int:
        return self._dim

class ModuleSpec:
    def __init__(self, name: str, module: nn.Module, in_dim: int, out_dim: int):
        self.name = name
        self.module = module
        self.in_dim = in_dim
        self.out_dim = out_dim

class QDMN(nn.Module):
    def __init__(self, initial_dim: int):
        super().__init__()
        self.core = CoreNode(initial_dim)
        self.initial_dim = initial_dim
        # registry: in_dim -> {symbol -> ModuleSpec}
        self._registry: Dict[int, Dict[str, ModuleSpec]] = {}

    def register_module(self, name: str, module: nn.Module, in_dim: int, out_dim: int,

```

```

        replace: bool=False):
    if in_dim not in self._registry:
        self._registry[in_dim] = {}
    if name in self._registry[in_dim] and not replace:
        raise KeyError(f"Module_{name}_{in_dim}_{already_registered}")
    self._registry[in_dim][name] = ModuleSpec(name, module, in_dim, out_dim)
    # register submodule for PyTorch tracking
    self.add_module(f"mod_{name}_{in_dim}", module)

def _find_module_for_step(self, step: str, cur_dim: int) -> ModuleSpec:
    if cur_dim not in self._registry:
        raise RuntimeError(f"No_modules_registered_for_current_dim_{cur_dim}")
    if step not in self._registry[cur_dim]:
        raise RuntimeError(f"Module_{step}_not_registered_for_dim_{cur_dim}")
    return self._registry[cur_dim][step]

def route(self, path: List[str], x: torch.Tensor, verbose: bool=False) -> torch.
Tensor:
    # x shape (... , feature_dim)
    if self.core._state is None: # allow core to accept this first input
        self.core.set(x)
    else:
        # require input last-dim to match initial dim for reproducibility
        if x.shape[-1] != self.initial_dim:
            # allow but warn
            pass
        self.core.set(x)

    if verbose:
        print(f"QDMN.route_start_dim={self.core.dim},_path='{'->'.join(path)}")
    for i, step in enumerate(path):
        cur_dim = self.core.dim
        spec = self._find_module_for_step(step, cur_dim)
        out = spec.module(self.core.get())
        if out.shape[-1] != spec.out_dim:
            raise RuntimeError(f"Module_{step}_returned_dim_{out.shape[-1]}_but_
                expected_{spec.out_dim}")
        self.core.set(out)
        if verbose:
            print(f"_{step}_{i+1}/{len(path)}:_{step}_{spec.in_dim}->{spec.out_dim}")
    if verbose:
        print("QDMN.route_finished,_final_dim=", self.core.dim)
    return self.core.get()

# -----
# Small module factories (examples)
# -----
class LinearBlock(nn.Module):
    def __init__(self, in_dim:int, out_dim:int, act:Optional[nn.Module]=nn.ReLU):
        super().__init__()
        self.lin = nn.Linear(in_dim, out_dim)
        self.act = act() if act is not None else nn.Identity()
    def forward(self, x: torch.Tensor) -> torch.Tensor:
        return self.act(self.lin(x))

```

```

class ConvSmall(nn.Module):
    """Simple conv backbone for tiny images; outputs flattened last-dim"""
    def __init__(self, in_channels:int, out_dim:int):
        super().__init__()
        # operate on (N, C, H, W) and output (N, out_dim) as last-dim feature
        self.net = nn.Sequential(
            nn.Conv2d(in_channels, 16, 3, padding=1), nn.ReLU(),
            nn.Conv2d(16, 32, 3, padding=1), nn.ReLU(),
            nn.AdaptiveAvgPool2d(1),
            nn.Flatten(),
            nn.Linear(32, out_dim),
            nn.ReLU()
        )
    def forward(self, x):
        return self.net(x)

class RNNBlock(nn.Module):
    """Small RNN block that consumes last-dim as sequence length; wrapper for demonstration only."""
    def __init__(self, seq_len:int, in_dim:int, out_dim:int, hidden:int=64):
        super().__init__()
        self.seq_len = seq_len
        self.rnn = nn.LSTM(in_dim, hidden, batch_first=True)
        self.head = nn.Linear(hidden, out_dim)
    def forward(self, x):
        # x shape (N, seq_len*in_dim) or (N, seq_len, in_dim)
        if x.dim()==2:
            x = x.view(x.size(0), self.seq_len, -1)
        out, _ = self.rnn(x)
        return self.head(out[:, -1, :])

# -----
# Baseline models
# -----
class BaselineMLP(nn.Module):
    def __init__(self, input_dim:int, hidden:int=128, out_dim:int=None):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(input_dim, hidden),
            nn.ReLU(),
            nn.Linear(hidden, out_dim if out_dim is not None else input_dim)
        )
    def forward(self, x): return self.net(x.view(x.size(0), -1))

class BaselineCNN(nn.Module):
    def __init__(self, in_channels:int, out_dim:int):
        super().__init__()
        self.net = ConvSmall(in_channels, out_dim)
    def forward(self, x): return self.net(x)

class BaselineRNN(nn.Module):
    def __init__(self, seq_len:int, in_dim:int, out_dim:int, hidden:int=64):
        super().__init__()

```

```

        self.net = RNNBlock(seq_len, in_dim, out_dim, hidden)
    def forward(self, x): return self.net(x)

# -----
# Toy datasets (paper use these or swap in real datasets)
# -----
def make_image_dataset(n_samples=1000, img_size=8):
    X = np.zeros((n_samples, 1, img_size, img_size), dtype=np.float32)
    y = np.zeros(n_samples, dtype=np.int64)
    for i in range(n_samples):
        cls = np.random.randint(0,2)
        y[i] = cls
        img = np.zeros((img_size, img_size), dtype=np.float32)
        if cls == 0:
            img[1:4,1:4] = 1.0
        else:
            img[4:7,4:7] = 1.0
        img += np.random.normal(0, 0.05, size=img.shape)
        X[i,0] = img
    return X, y

def make_sine_regression(n_samples=1000):
    X = np.random.uniform(-1, 1, size=(n_samples,1)).astype(np.float32)
    y = (np.sin(2*np.pi*X[:,0]) + np.random.normal(0, 0.1, size=n_samples)).astype(np.
        float32).reshape(-1,1)
    return X, y

def make_sequence_parity(n_samples=1000, seq_len=10):
    X = np.random.randint(0,2,size=(n_samples, seq_len)).astype(np.float32)
    y = (X.sum(axis=1) % 2).astype(np.int64)
    return X, y

# -----
# Training / evaluation helpers
# -----
def count_params(m:nn.Module):
    return sum(p.numel() for p in m.parameters() if p.requires_grad)

def train_baseline(model, train_loader, val_loader, epochs, lr, device, regression=False,
    verbose=False):
    model.to(device)
    opt = optim.Adam(model.parameters(), lr=lr)
    loss_fn = nn.MSELoss() if regression else nn.CrossEntropyLoss()
    for ep in range(epochs):
        model.train()
        for xb, yb in train_loader:
            xb = xb.to(device); yb = yb.to(device)
            opt.zero_grad()
            out = model(xb)
            loss = loss_fn(out, yb)
            loss.backward(); opt.step()

```

```

        if verbose and (ep+1) % 20 == 0:
            print(f"epoch_{ep+1}/{epochs} train_loss_{loss.item():.6f}")
# eval
model.eval()
with torch.no_grad():
    total_loss = 0.0; total_count=0
    if regression:
        for xb,yb in val_loader:
            xb=xb.to(device); yb=yb.to(device)
            pred = model(xb)
            total_loss += ((pred-yb)**2).sum().item()
            total_count += yb.numel()
        return total_loss / total_count
    else:
        correct=0; total=0
        for xb,yb in val_loader:
            xb=xb.to(device); yb=yb.to(device)
            pred = model(xb).argmax(dim=1)
            correct += (pred==yb).sum().item(); total += xb.size(0)
        return correct/total

def train_qdmn(qdmn:QDMN, route:List[str], train_loader, val_loader, epochs, lr, device,
               regression=False, verbose=False):
    qdmn.to(device)
    opt = optim.Adam(qdmn.parameters(), lr=lr)
    loss_fn = nn.MSELoss() if regression else nn.CrossEntropyLoss()
    for ep in range(epochs):
        qdmn.train()
        for xb,yb in train_loader:
            xb=xb.to(device); yb=yb.to(device)
            opt.zero_grad()
            out = qdmn.route(route, xb, verbose=False)
            loss = loss_fn(out, yb)
            loss.backward(); opt.step()
        if verbose and (ep+1) % 20 == 0:
            print(f"epoch_{ep+1}/{epochs} train_loss_{loss.item():.6f}")
# eval
qdmn.eval()
with torch.no_grad():
    if regression:
        tot=0.0; cnt=0
        for xb,yb in val_loader:
            xb=xb.to(device); yb=yb.to(device)
            pred = qdmn.route(route, xb, verbose=False)
            tot += ((pred-yb)**2).sum().item(); cnt += yb.numel()
        return tot / cnt
    else:
        correct=0; tot=0
        for xb,yb in val_loader:
            xb=xb.to(device); yb=yb.to(device)
            pred = qdmn.route(route, xb, verbose=False).argmax(dim=1)
            correct += (pred==yb).sum().item(); tot += xb.size(0)
        return correct / tot

```

```

def measure_inference_time(model_callable, sample, runs=100, device=DEVICE):
    # model_callable: function that runs forward pass on sample (already on correct
    # device)
    # returns avg ms per run
    with torch.no_grad():
        # warmup
        for _ in range(5):
            model_callable()
        t0 = time.time()
        for _ in range(runs):
            model_callable()
        t_total = time.time() - t0
    return (t_total / runs) * 1000.0

# -----
# Benchmark harness
# -----
def run_benchmarks(print_plots=False):
    random.seed(SEED); np.random.seed(SEED); torch.manual_seed(SEED)

    results = []

    # ---- Task A: Image classification (tiny 8x8 images) ----
    X, y = make_image_dataset(n_samples=1200, img_size=8)
    Xtr, Xval, ytr, yval = train_test_split(X, y, test_size=0.2, random_state=SEED)
    # Baseline CNN (operates on (N,1,8,8)), Baseline MLP (flattened), QDMN works on
    # flattened vector
    bs = BATCH
    train_cnn = DataLoader(TensorDataset(torch.tensor(Xtr), torch.tensor(ytr)),
                           batch_size=bs, shuffle=True)
    val_cnn = DataLoader(TensorDataset(torch.tensor(Xval), torch.tensor(yval)),
                         batch_size=bs)

    # Baseline CNN
    cnn = BaselineCNN(in_channels=1, out_dim=16) # produce 16-d features
    # add final classifier head
    cnn_head = nn.Linear(16, 2)
    cnn_full = nn.Sequential(cnn, cnn_head)
    acc_cnn = train_baseline(cnn_full, train_cnn, val_cnn, epochs=EPOCHS_SMALL, lr=LR,
                             device=DEVICE, regression=False)

    # Baseline MLP (flatten)
    Xtr_flat = Xtr.reshape(len(Xtr), -1); Xval_flat = Xval.reshape(len(Xval), -1)
    train_mlp = DataLoader(TensorDataset(torch.tensor(Xtr_flat), torch.tensor(ytr)),
                           batch_size=bs, shuffle=True)
    val_mlp = DataLoader(TensorDataset(torch.tensor(Xval_flat), torch.tensor(yval)),
                         batch_size=bs)
    mlp = BaselineMLP(input_dim=Xtr_flat.shape[1], hidden=128, out_dim=2)
    acc_mlp = train_baseline(mlp, train_mlp, val_mlp, epochs=EPOCHS_SMALL, lr=LR, device=
                             DEVICE, regression=False)

    # QDMN variant: register modules on flattened dim
    DIM = Xtr_flat.shape[1]

```

```

qdmn_img = QDMN(initial_dim=DIM)
# F, B preserve dim; U compress to 32, D reconstruct; H head maps to logits
qdmn_img.register_module('F', LinearBlock(DIM, DIM, act=nn.ReLU), in_dim=DIM, out_dim=DIM)
qdmn_img.register_module('B', LinearBlock(DIM, DIM, act=nn.Tanh), in_dim=DIM, out_dim=DIM)
qdmn_img.register_module('U', LinearBlock(DIM, 32, act=nn.ReLU), in_dim=DIM, out_dim=32)
qdmn_img.register_module('F', LinearBlock(32, 32, act=nn.ReLU), in_dim=32, out_dim=32)
qdmn_img.register_module('D', LinearBlock(32, DIM, act=nn.ReLU), in_dim=32, out_dim=DIM)
qdmn_img.register_module('H', nn.Linear(DIM, 2), in_dim=DIM, out_dim=2) # classifier head
# Train QDMN on route ['F','H'] (short) and evaluate also on ['U','D','H']
acc_q_short = train_qdmn(qdmn_img, ['F','H'], train_mlp, val_mlp, epochs=EPOCHS_SMALL,
    lr=LR, device=DEVICE, regression=False)
acc_q_comp = train_qdmn(qdmn_img, ['U','D','H'], train_mlp, val_mlp, epochs=EPOCHS_SMALL,
    lr=LR, device=DEVICE, regression=False)

results.append({
    'task': 'Image_(8x8)_classification',
    'metric': 'accuracy',
    'CNN': round(acc_cnn,4),
    'MLP': round(acc_mlp,4),
    'QDMN_[F,H]': round(acc_q_short,4),
    'QDMN_[U,D,H]': round(acc_q_comp,4),
    'notes': 'Flattened_for_QDMN'
})
print("Image_task_done.")

# ---- Task B: Regression (sine) ----
Xr, yr = make_sine_regression(n_samples=1200)
Xtr, Xval, ytr, yval = train_test_split(Xr, yr, test_size=0.2, random_state=SEED)
train_r = DataLoader(TensorDataset(torch.tensor(Xtr), torch.tensor(ytr)), batch_size=bs, shuffle=True)
val_r = DataLoader(TensorDataset(torch.tensor(Xval), torch.tensor(yval)), batch_size=bs)

mlp_reg = BaselineMLP(input_dim=1, hidden=128, out_dim=1)
mse_mlp_reg = train_baseline(mlp_reg, train_r, val_r, epochs=EPOCHS_LARGE, lr=LR,
    device=DEVICE, regression=True)

# QDMN regressor: expand -> process -> decode
q_reg = QDMN(initial_dim=1)
q_reg.register_module('F', LinearBlock(1, 32, act=nn.ReLU), in_dim=1, out_dim=32)
q_reg.register_module('B', LinearBlock(32, 32, act=nn.Tanh), in_dim=32, out_dim=32)
q_reg.register_module('D', nn.Linear(32, 1), in_dim=32, out_dim=1)
mse_q = train_qdmn(q_reg, ['F','B','D'], train_r, val_r, epochs=EPOCHS_LARGE, lr=LR,
    device=DEVICE, regression=True)

results.append({
    'task': 'Regression_(sine)',
    'metric': 'MSE_(lower_better)',

```

```

        'MLP': round(mse_mlp_reg,6),
        'QDMN_[F,B,D]': round(mse_q,6),
        'notes': '1D_expand/process/decode'
    })
    print("Regression_task_done.")

# ---- Task C: Sequence parity ----
seq_len = 10
Xs, ys = make_sequence_parity(n_samples=1200, seq_len=seq_len)
Xtr, Xval, ytr, yval = train_test_split(Xs, ys, test_size=0.2, random_state=SEED)
# Baselines:
train_seq_mlp = DataLoader(TensorDataset(torch.tensor(Xtr.reshape(len(Xtr), -1)),
    torch.tensor(ytr)), batch_size=bs, shuffle=True)
val_seq_mlp = DataLoader(TensorDataset(torch.tensor(Xval.reshape(len(Xval), -1)),
    torch.tensor(yval)), batch_size=bs)
train_seq_rnn = DataLoader(TensorDataset(torch.tensor(Xtr), torch.tensor(ytr)),
    batch_size=bs, shuffle=True)
val_seq_rnn = DataLoader(TensorDataset(torch.tensor(Xval), torch.tensor(yval)),
    batch_size=bs)

mlp_seq = BaselineMLP(input_dim=seq_len, hidden=128, out_dim=2)
acc_mlp_seq = train_baseline(mlp_seq, train_seq_mlp, val_seq_mlp, epochs=EPOCHS_LARGE
    , lr=LR, device=DEVICE, regression=False)

rnn_seq = BaselineRNN(seq_len=seq_len, in_dim=1, out_dim=2, hidden=64)
acc_rnn_seq = train_baseline(rnn_seq, train_seq_rnn, val_seq_rnn, epochs=EPOCHS_LARGE
    , lr=LR, device=DEVICE, regression=False)

# QDMN: treat flattened seq as vector, use compress/process/decode route
q_seq = QDMN(initial_dim=seq_len)
q_seq.register_module('F', LinearBlock(seq_len, 32, act=nn.ReLU), in_dim=seq_len,
    out_dim=32)
q_seq.register_module('B', LinearBlock(32, 32, act=nn.Tanh), in_dim=32, out_dim=32)
q_seq.register_module('D', LinearBlock(32, seq_len, act=nn.ReLU), in_dim=32, out_dim=
    seq_len)
q_seq.register_module('H', nn.Linear(seq_len, 2), in_dim=seq_len, out_dim=2)
acc_q_seq = train_qdmn(q_seq, ['F','B','D','H'], train_seq_mlp, val_seq_mlp, epochs=
    EPOCHS_LARGE, lr=LR, device=DEVICE, regression=False)

results.append({
    'task': 'Sequence_parity_(len=10)',
    'metric': 'accuracy',
    'MLP': round(acc_mlp_seq,4),
    'RNN': round(acc_rnn_seq,4),
    'QDMN_[F,B,D,H]': round(acc_q_seq,4)
})
print("Sequence_parity_task_done.")

# ---- Novelty / Route-switching test (single clear example for the paper) ----
# Train QDMN on short route ['F','B'] for regression and test on longer routes
print("\nRunning_Route-Switching_(Novelty)_test...")
# reuse q_reg (regression qdmn) but make a fresh instance to be clean
q_novel = QDMN(initial_dim=1)
q_novel.register_module('F', LinearBlock(1, 32, act=nn.ReLU), in_dim=1, out_dim=32)

```

```

q_novel.register_module('B', LinearBlock(32, 32, act=nn.Tanh), in_dim=32, out_dim=32)
q_novel.register_module('D', nn.Linear(32, 1), in_dim=32, out_dim=1)

# prepare regression loaders
Xr, Yr = make_sine_regression(n_samples=2048)
Xtr, Xv, Ytr, Yv = train_test_split(Xr, Yr, test_size=0.2, random_state=SEED)
train_loader_n = DataLoader(TensorDataset(torch.tensor(Xtr), torch.tensor(Ytr)),
                             batch_size=bs, shuffle=True)
val_loader_n = DataLoader(TensorDataset(torch.tensor(Xv), torch.tensor(Yv)),
                           batch_size=bs)

# Train on short route
train_qdmn(q_novel, ['F', 'B', 'D'], train_loader_n, val_loader_n, epochs=EPOCHS_LARGE,
            lr=LR, device=DEVICE, regression=True)
# Evaluate on short route
mse_short = train_qdmn(q_novel, ['F', 'B', 'D'], train_loader_n, val_loader_n, epochs
                        =0, lr=LR, device=DEVICE, regression=True)
# Evaluate on longer unseen route (more passes)
# create longer route by repeating B (refinements) and finishing with D
long_route = ['F', 'B', 'B', 'B', 'D']
mse_long = train_qdmn(q_novel, long_route, train_loader_n, val_loader_n, epochs=0, lr
                      =LR, device=DEVICE, regression=True)
novelty_flag = mse_long < mse_short

results.append({
    'task': 'Novelty: route-switching (regression)',
    'metric': 'MSE (short vs long)',
    'mse_short': round(mse_short, 6),
    'mse_long': round(mse_long, 6),
    'novelty_improved': bool(novelty_flag)
})
print("Novelty test done. improved?:", novelty_flag)

# Consolidate results
df = pd.DataFrame(results)
print("\n=== Benchmark Summary ===")
print(df.to_markdown(index=False))

# Save CSV for paper artifact
os.makedirs("qdmn_artifacts", exist_ok=True)
df.to_csv("qdmn_artifacts/benchmark_summary.csv", index=False)
print("Saved results to qdmn_artifacts/benchmark_summary.csv")

# optional: show a small plot for the novelty test
try:
    plt.figure(figsize=(4,3))
    plt.bar(['short', 'long'], [mse_short, mse_long])
    plt.title("Novelty test (MSE): short vs long route")
    plt.ylabel("MSE (lower better)")
    plt.show()
except Exception:
    pass

return df

```

```
# Run
if __name__ == "__main__":
    df = run_benchmarks()
    print("\nDone. Increase EPOCHS_* and model sizes to make stronger claims for your paper.")
```

B Depth Generalization Benchmark (Stable vs Naive)

The following script implements a focused experiment comparing naive depth-1 training to multi-depth training using a gated residual refiner and reports MSE across inference depths.

```
# =====
# QDMN FINAL PAPER BENCHMARK
# Stable vs Unstable Depth Generalization
# =====

import os, random, time
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt

# =====
# Global config
# =====
SEED = 42
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
OUTDIR = "qdmn_final_artifacts"
os.makedirs(OUTDIR, exist_ok=True)

EPOCHS = 1000
LR = 1e-3
DEPTHS = [1, 2, 4, 8]

print("Device:", DEVICE)
print("Artifacts:", OUTDIR)

# =====
# Reproducibility
# =====
def set_seed(seed):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)

set_seed(SEED)

def make_task(n: int = 512):
    x = np.random.uniform(-1, 1, (n, 1)).astype(np.float32)
```

```

    y = np.sin(2 * np.pi * x).astype(np.float32)
    return torch.tensor(x), torch.tensor(y)

X, Y = make_task()
X, Y = X.to(DEVICE), Y.to(DEVICE)

class MLP(nn.Module):
    def __init__(self):
        super().__init__()
        self.net = nn.Sequential(nn.Linear(1, 64), nn.ReLU(), nn.Linear(64, 1))

    def forward(self, x):
        return self.net(x)

class ContractiveRefiner(nn.Module):
    def __init__(self, dim: int):
        super().__init__()
        self.W = nn.utils.parametrizations.spectral_norm(nn.Linear(dim, dim))
        self.gate = nn.Linear(dim, 1)
        self.alpha = 0.1

    def forward(self, x):
        g = torch.sigmoid(self.gate(x))
        delta = torch.tanh(self.W(x))
        return x + self.alpha * g * delta

class QDMN(nn.Module):
    def __init__(self):
        super().__init__()
        self.F = nn.Linear(1, 32)
        self.B = ContractiveRefiner(32)
        self.D = nn.Linear(32, 1)

    def forward(self, x, depth: int):
        s = torch.relu(self.F(x))
        for _ in range(depth):
            s = self.B(s)
        return self.D(s)

# =====
# Toy regression task
# =====
def make_task(n=512):
    x = np.random.uniform(-1, 1, (n, 1)).astype(np.float32)
    y = np.sin(2 * np.pi * x)
    return torch.tensor(x), torch.tensor(y)

X, Y = make_task()
X, Y = X.to(DEVICE), Y.to(DEVICE)

```

```

# =====
# Baseline MLP
# =====
class MLP(nn.Module):
    def __init__(self):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(1, 64),
            nn.ReLU(),
            nn.Linear(64, 1)
        )
    def forward(self, x):
        return self.net(x)

# =====
# QDMN Components
# =====
class ContractiveRefiner(nn.Module):
    """
    Spectrally-normalized, gated residual refiner
    Guarantees bounded iteration
    """
    def __init__(self, dim):
        super().__init__()
        self.W = nn.utils.parametrizations.spectral_norm(
            nn.Linear(dim, dim)
        )
        self.gate = nn.Linear(dim, 1)
        self.alpha = 0.1

    def forward(self, x):
        g = torch.sigmoid(self.gate(x))
        delta = torch.tanh(self.W(x))
        return x + self.alpha * g * delta

class QDMN(nn.Module):
    def __init__(self):
        super().__init__()
        self.F = nn.Linear(1, 32)
        self.B = ContractiveRefiner(32)
        self.D = nn.Linear(32, 1)

    def forward(self, x, depth):
        s = torch.relu(self.F(x))
        for _ in range(depth):
            s = self.B(s)
        return self.D(s)

# =====
# Training helpers
# =====
def train(model, loss_fn, optimizer, depths=[1]):
    for _ in range(EPOCHS):
        optimizer.zero_grad()

```

```

        loss = 0
        for d in depths:
            loss += loss_fn(model(X, d), Y)
        loss /= len(depths)
        loss.backward()
        optimizer.step()

def eval_depths(model, depths):
    with torch.no_grad():
        return {
            d: nn.MSELoss()(model(X, d), Y).item()
            for d in depths
        }

# =====
# 1. Baseline MLP
# =====
mlp = MLP().to(DEVICE)
opt = optim.Adam(mlp.parameters(), LR)
for _ in range(EPOCHS):
    opt.zero_grad()
    loss = nn.MSELoss()(mlp(X), Y)
    loss.backward()
    opt.step()

mlp_loss = nn.MSELoss()(mlp(X), Y).item()

# =====
# 2. QDMN Naive (depth=1 training)
# =====
q_naive = QDMN().to(DEVICE)
opt = optim.Adam(q_naive.parameters(), LR)
train(q_naive, nn.MSELoss(), opt, depths=[1])
naive_results = eval_depths(q_naive, DEPTHS)

# =====
# 3. QDMN Stable (multi-depth training)
# =====
q_stable = QDMN().to(DEVICE)
opt = optim.Adam(q_stable.parameters(), LR)
train(q_stable, nn.MSELoss(), opt, depths=[1, 2, 4])
stable_results = eval_depths(q_stable, DEPTHS)

# =====
# Results Table
# =====
df = pd.DataFrame({
    "Depth": DEPTHS,
    "QDMN_Naive": [naive_results[d] for d in DEPTHS],
    "QDMN_Stable": [stable_results[d] for d in DEPTHS]
})

print("\n===FINAL_RESULTS_TABLE===")
print(df)

```

```

# =====
# Plot
# =====
plt.figure(figsize=(6,4))
plt.plot(DEPTHS, df["QDMN_Naive"], "o--", label="QDMN_Naive")
plt.plot(DEPTHS, df["QDMN_Stable"], "o-", label="QDMN_Stable")
plt.axhline(mlp_loss, color="gray", linestyle=":", label="Baseline_MLP")
plt.xscale("log", base=2)
plt.xlabel("Inference_Depth")
plt.ylabel("MSE")
plt.title("Depth_Generalization_Behavior")
plt.legend()
plt.grid(True, alpha=0.3)
plt.savefig(f"{OUTDIR}/depth_generalization.png", dpi=200)
plt.show()

# =====
# Interpretation Print
# =====
print("\n===_INTERPRETATION_===")
print("_Naive_QDMN_diverges_with_depth_expected")
print("_Stable_QDMN_remains_bounded_framework+_training_dependent")
print("_Depth_generalization_is_conditional,_not_automatic")

print("\nDone._All_artifacts_saved_to:", OUTDIR)

```

C Semantic Compositionality Experiment Script

Script for the [F, U, D, Bⁿ, H] noise-filtering demo (Figure 2).

```

# QDMN semantic compositionality experiment
# Output: qdmn_final_artifacts/qdmn_composition_test.png

# =====
# QDMN MASTER EXPERIMENT: Semantic Compositionality [F, U, D, Bn]
# Purpose: Proves the framework can compose directions to solve
# complex noise-filtering tasks via "Route-Based Logic."
# =====

import os
import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt

OUTDIR = "qdmn_final_artifacts"
os.makedirs(OUTDIR, exist_ok=True)

# --- 1. The Quad-Directional Registry ---
class ContractiveB(nn.Module):
    def __init__(self, dim):
        super().__init__()

```

```

        self.net = nn.utils.parametrizations.spectral_norm(nn.Linear(dim, dim))
        self.gate = nn.Linear(dim, 1)
    def forward(self, x):
        g = torch.sigmoid(self.gate(x))
        return x + 0.1 * g * torch.tanh(self.net(x))

class QDMN_Master(nn.Module):
    def __init__(self):
        super().__init__()
        # Forward: Initial Projection
        self.F = nn.Linear(1, 64)
        # Up: Abstraction (Dimensionality Reduction)
        self.U = nn.Linear(64, 16)
        # Down: Reconstruction (Dimensionality Expansion)
        self.D = nn.Linear(16, 64)
        # Backward: Iterative Refinement
        self.B = ContractiveB(64)
        # Head: Output mapping
        self.H = nn.Linear(64, 1)

    def route(self, path, x):
        s = x
        for step in path:
            if step == 'F': s = torch.relu(self.F(s))
            elif step == 'U': s = torch.relu(self.U(s))
            elif step == 'D': s = torch.relu(self.D(s))
            elif step == 'B': s = self.B(s)
            elif step == 'H': s = self.H(s)
        return s

# --- 2. Data: Noisy Sine Wave ---
x = torch.linspace(-1, 1, 500).view(-1, 1)
clean_y = torch.sin(3 * x)
noisy_y = clean_y + 0.3 * torch.randn_like(clean_y) # 30% noise

# --- 3. Training the Composition ---
model = QDMN_Master()
optimizer = optim.Adam(model.parameters(), lr=0.01)

print("Training Master Route: [F->U->D->B]...")
for i in range(1000):
    optimizer.zero_grad()
    # Train the model to clean noise using a specific semantic route
    pred = model.route(['F', 'U', 'D', 'B', 'H'], noisy_y)
    loss = nn.MSELoss()(pred, clean_y)
    loss.backward()
    optimizer.step()

# --- 4. Inference: The "Power of the Route" ---
with torch.no_grad():
    # Route 1: Fast (No refinement)
    y_fast = model.route(['F', 'H'], noisy_y)
    # Route 2: Deep (Full Semantic Composition)
    y_full = model.route(['F', 'U', 'D', 'B', 'B', 'B', 'H'], noisy_y)

```

```

# --- 5. Visualization for Paper ---
plt.figure(figsize=(10, 5))
plt.scatter(x, noisy_y, alpha=0.2, label="Noisy_Input", color='gray')
plt.plot(x, clean_y, 'k--', label="Ground_Truth", lw=2)
plt.plot(x, y_fast, label="Short_Route_[F,H]", color='red')
plt.plot(x, y_full, label="Master_Route_[F,U,D,B^3,H]", color='blue', lw=3)
plt.title("QDMN_Semantic_Compositionality:_Noise_Filtering_via_Programmable_Routes")
plt.legend()
out_path = os.path.join(OUTDIR, "qdmn_composition_test.png")
plt.savefig(out_path, dpi=300)
plt.show()

print(f"Experiment_complete._Figure_saved_to_{out_path}")

```

D Multi-seed Paper-Ready Experiment Script

Extended multi-seed experiment script (ablation, depth sweeps, and CSV outputs). This is the paper-ready version provided by the author.

```

# qdmn_paper_ready_full_fixed2.py
# Single-file reproducible experiments for QDMN (paper-ready)
# Requirements: torch, numpy, pandas, matplotlib, scikit-learn, (optional: scipy)

import os, time, random, math
from typing import List, Dict, Tuple

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import torch
import torch.nn as nn
import torch.optim as optim
from sklearn.model_selection import train_test_split

# optional: scipy for stats
try:
    from scipy import stats
    SCIPY_OK = True
except Exception:
    SCIPY_OK = False

# -----
# Config
# -----
SEEDS = [42, 123, 999]
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
OUTDIR = "qdmn_paper_artifacts"
os.makedirs(OUTDIR, exist_ok=True)

N_SAMPLES = 2048
BATCH = 128

```

```

EPOCHS_BASELINE = 120
EPOCHS_QDMN = 240
LR = 1e-3

EVAL Depths = [1, 2, 4, 8, 16, 32]

print("Device:", DEVICE)
print("Artifacts_dir:", OUTDIR)
print("scipy_available:", SCIPY_OK)

# -----
# Utilities
# -----
def set_seed(seed: int):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    if DEVICE.type == "cuda":
        torch.cuda.manual_seed_all(seed)

def mean_ci(a: List[float], alpha=0.05):
    arr = np.array(a)
    m = float(arr.mean())
    if len(arr) <= 1:
        return m, (0.0, 0.0)
    s = float(arr.std(ddof=1))
    n = len(arr)
    se = s / math.sqrt(n)
    z = float(stats.t.ppf(1 - alpha / 2, df=n - 1)) if SCIPY_OK else 1.96
    return m, (m - z * se, m + z * se)

def paired_ttest(a: List[float], b: List[float]) -> Tuple[float, float]:
    if SCIPY_OK:
        t, p = stats.ttest_rel(a, b)
        return float(t), float(p)
    return float("nan"), float("nan")

def measure_inference_time(fn, runs=50):
    with torch.no_grad():
        for _ in range(5):
            fn()
        t0 = time.time()
        for _ in range(runs):
            fn()
        t_total = time.time() - t0
    return (t_total / runs) * 1000.0

# -----
# Dataset
# -----
def make_sine(n=N_SAMPLES, noise_std=0.1):

```

```

x = np.random.uniform(-1.0, 1.0, size=(n, 1)).astype(np.float32)
y = (np.sin(2 * np.pi * x) + np.random.normal(0, noise_std, size=x.shape)).astype(np.
    float32)
return x, y

# -----
# Models
# -----
class BaselineMLP(nn.Module):
    def __init__(self, in_dim=1, hidden=64):
        super().__init__()
        self.net = nn.Sequential(nn.Linear(in_dim, hidden), nn.ReLU(), nn.Linear(hidden,
            1))

    def forward(self, x):
        return self.net(x.float())

class ContractiveRefiner(nn.Module):
    def __init__(self, dim):
        super().__init__()
        self.linear = nn.utils.spectral_norm(nn.Linear(dim, dim))
        self.gate = nn.Linear(dim, 1)
        self.alpha = 0.08

    def forward(self, x):
        g = torch.sigmoid(self.gate(x))
        delta = torch.tanh(self.linear(x))
        return x + self.alpha * g * delta

class PlainRefiner(nn.Module):
    def __init__(self, dim):
        super().__init__()
        self.net = nn.Sequential(nn.Linear(dim, dim), nn.Tanh())
        self.alpha = 0.1

    def forward(self, x):
        return x + self.alpha * self.net(x)

class QDMNModel(nn.Module):
    def __init__(self, refiner: nn.Module, hidden=32):
        super().__init__()
        self.F = nn.Linear(1, hidden)
        self.B = refiner
        self.D = nn.Linear(hidden, 1)

    def forward(self, x, depth=1):
        s = torch.relu(self.F(x.float()))
        for _ in range(depth):
            s = self.B(s)
        return self.D(s)

```

```

# -----
# Fitters
# -----
def fit_mse(model, train_loader, val_loader, epochs, lr, device):
    model.to(device)
    opt = optim.Adam(model.parameters(), lr=lr)
    loss_fn = nn.MSELoss()
    for _ in range(epochs):
        model.train()
        for xb, yb in train_loader:
            xb, yb = xb.to(device), yb.to(device)
            opt.zero_grad()
            loss = loss_fn(model(xb), yb)
            loss.backward()
            opt.step()

    model.eval()
    with torch.no_grad():
        tot = 0.0
        cnt = 0
        for xb, yb in val_loader:
            xb, yb = xb.to(device), yb.to(device)
            pred = model(xb)
            tot += ((pred - yb) ** 2).sum().item()
            cnt += yb.numel()
    return tot / max(cnt, 1)

def fit_qdmn(model: QDMNModel, train_loader, val_loader, depths_for_training, epochs, lr,
             device):
    model.to(device)
    opt = optim.Adam(model.parameters(), lr=lr)
    loss_fn = nn.MSELoss()

    for _ in range(epochs):
        model.train()
        for xb, yb in train_loader:
            xb, yb = xb.to(device), yb.to(device)
            opt.zero_grad()
            loss = 0.0
            for d in depths_for_training:
                loss = loss + loss_fn(model(xb, depth=d), yb)
            loss = loss / float(len(depths_for_training))
            loss.backward()
            opt.step()

    model.eval()
    with torch.no_grad():
        tot = 0.0
        cnt = 0
        for xb, yb in val_loader:
            xb, yb = xb.to(device), yb.to(device)
            pred = model(xb, depth=depths_for_training[0])
            tot += ((pred - yb) ** 2).sum().item()

```

```

        cnt += yb.numel()
    return tot / max(cnt, 1)

# -----
# Run per seed
# -----
def run_seed(seed: int):
    set_seed(seed)
    X_all, Y_all = make_sine(N_SAMPLES)
    X_trainval, X_test, y_trainval, y_test = train_test_split(X_all, Y_all, test_size
        =0.1, random_state=seed)
    X_train, X_val, y_train, y_val = train_test_split(X_trainval, y_trainval, test_size
        =0.1111111, random_state=seed)

    Xtr = torch.tensor(X_train)
    ytr = torch.tensor(y_train)
    Xval = torch.tensor(X_val)
    yval = torch.tensor(y_val)

    train_loader = torch.utils.data.DataLoader(torch.utils.data.TensorDataset(Xtr, ytr),
        batch_size=BATCH, shuffle=True)
    val_loader = torch.utils.data.DataLoader(torch.utils.data.TensorDataset(Xval, yval),
        batch_size=BATCH, shuffle=False)

    mlp = BaselineMLP(hidden=64)
    mlp_val = fit_mse(mlp, train_loader, val_loader, epochs=EPOCHS_BASELINE, lr=LR,
        device=DEVICE)

    q_naive = QDMNModel(refiner=PlainRefiner(32), hidden=32)
    qn_val = fit_qdmn(q_naive, train_loader, val_loader, depths_for_training=[1], epochs=
        EPOCHS_QDMN, lr=LR, device=DEVICE)

    q_stable = QDMNModel(refiner=ContractiveRefiner(32), hidden=32)
    qs_val = fit_qdmn(q_stable, train_loader, val_loader, depths_for_training=[1, 2, 4],
        epochs=EPOCHS_QDMN, lr=LR, device=DEVICE)

    # Depth sweep on validation set
    Xval_d = Xval.to(DEVICE)
    yval_d = yval.to(DEVICE)
    loss_fn = nn.MSELoss()
    depth_rows = []
    with torch.no_grad():
        depth_rows.append({"model": "MLP", "depth": 1, "mse": float(loss_fn(mlp(Xval_d),
            yval_d).item())})
    for d in EVAL_DEPTHS:
        depth_rows.append({"model": "QDMN_naive", "depth": d, "mse": float(loss_fn(
            q_naive(Xval_d, depth=d), yval_d).item())})
        depth_rows.append({"model": "QDMN_stable", "depth": d, "mse": float(loss_fn(
            q_stable(Xval_d, depth=d), yval_d).item())})

    return {"seed": seed, "mlp_val": mlp_val, "qnaive_val": qn_val, "qstable_val": qs_val
        , "depth_rows": depth_rows}

```

```

if __name__ == "__main__":
    rows = []
    depth_rows = []

    for s in SEEDS:
        print("Running seed", s)
        res = run_seed(s)
        rows.append({k: res[k] for k in ["seed", "mlp_val", "qnaive_val", "qstable_val"]})
        for d in res["depth_rows"]:
            depth_rows.append({"seed": s, **d})

    df = pd.DataFrame(rows)
    df_depths = pd.DataFrame(depth_rows)

    df.to_csv(os.path.join(OUTDIR, "per_seed_summary.csv"), index=False)
    df_depths.to_csv(os.path.join(OUTDIR, "depths_summary.csv"), index=False)

    print("\nSaved:")
    print("└-", os.path.join(OUTDIR, "per_seed_summary.csv"))
    print("└-", os.path.join(OUTDIR, "depths_summary.csv"))

```